

time	rows	estimation	cost
#1	HashAggr...		
#2	↳ Nested L...		
#3	↳ Seq Sc...		
#4	↳ Memo...		
#5	↳ Ind...		

#1 HashAggregate🕒 \$ 🗨

by departments.department_name

HashAggregate Node groups records together based on a GROUP BY or aggregate function (like sum()). Hash Aggregate uses a hash to first organize the records by a key.

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.284ms | 28%

☰ Rows: 11 (Planned: 200) | ⬆ over estimated by 18 ×

\$ Cost: 5.7 (Total: 45.4)

#2 Nested Loop🕒 \$

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.454ms | 45%

☰ Rows: 741 (Planned: 741)

\$ Cost: 21.1 (Total: 39.7)

#3 Seq Scan🕒 \$

on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.233ms | 23%

☰ Rows: 741 (Planned: 741)

\$ Cost: 18.4 (Total: 18.4)

#4 Memoize

Memoize Node is used to cache the results of the inner side of a nested loop. It avoids executing underlying nodes when the results for the current parameters are already in the cache.

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0ms | 0%

☰ Rows: ~741 (Planned: ~741)

\$ Cost: 0.01 (Total: 0.24)

🔄 Loops: 741

#5 Index Scan

on departments
using departments_pkey

Index Scan Node finds relevant records based on an Index. Index Scans perform 2 read operations: one to read the index and another to read the actual value from the table.

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.033ms | 3%

☰ Rows: ~11 (Planned: ~11)

\$ Cost: 0.23 (Total: 0.23)

🔄 Loops: 11

time	rows	estimation	cost
#1	HashAggr...		
#2	└ Hash Join		
#3	└└ Seq Sc...		
#4	└└ Hash		
#5	└└└ Seq...		

#1 HashAggregate⌚ \$

by departments.department_name

HashAggregate Node groups records together based on a GROUP BY or aggregate function (like sum()). Hash Aggregate uses a hash to first organize the records by a key.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.311ms | 33%

☰ Rows: 11 (Planned: 11)

\$ Cost: 3.81 (Total: 26.2)

#2 Hash Join⌚ \$

on employees.department_id = departments.department_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.452ms | 48%

☰ Rows: 741 (Planned: 741)

\$ Cost: 2.83 (Total: 22.4)

#3 Seq Scan⌚ \$

on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.169ms | 18%

☰ Rows: 741 (Planned: 741)

\$ Cost: 18.4 (Total: 18.4)

#4 Hash

Hash Node generates a hash table from the records in the input recordset. Hash is used by Hash Join.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.007ms | 1%

☰ Rows: 11 (Planned: 11)

#5 Seq Scan

on departments

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.009ms | 1%

☰ Rows: 11 (Planned: 11)

\$ Cost: 1.11 (Total: 1.11)

time	rows	estimation	cost
#1	HashAggr...		
#2	└ Hash Join		
#3	└└ Seq Sc...		
#4	└└ Hash		
#5	└└└ Seq...		

#1 HashAggregate

by locations.city

HashAggregate Node groups records together based on a GROUP BY or aggregate function (like sum()). Hash Aggregate uses a hash to first organize the records by a key.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.011ms | 16%

☰ Rows: 7 (Planned: 11)

\$ Cost: 0.16 (Total: 14.8)

#2 Hash Join

on locations.location_id = departments.location_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.015ms | 22%

☰ Rows: 11 (Planned: 11)

\$ Cost: 1.15 (Total: 14.7)

#3 Seq Scan

on locations

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.007ms | 10%

☰ Rows: 7 (Planned: 240) | ⬆ over estimated by 34 ×

\$ Cost: 12.4 (Total: 12.4)

#4 Hash

Hash Node generates a hash table from the records in the input recordset. Hash is used by Hash Join.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.009ms | 13%

☰ Rows: 11 (Planned: 11)

#5 Seq Scan

on departments

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.026ms | 38%

☰ Rows: 11 (Planned: 11)

\$ Cost: 1.11 (Total: 1.11)

time	rows	estimation	cost
#1	HashAggr...		
#2	└ Hash Join		
#3	└└ Seq Sc...		
#4	└└ Hash		
#5	└└└ Seq...		

#1 HashAggregate

by locations.city

HashAggregate Node groups records together based on a GROUP BY or aggregate function (like sum()). Hash Aggregate uses a hash to first organize the records by a key.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.01ms | 9%

☰ Rows: 7 (Planned: 7)

\$ Cost: 0.12 (Total: 2.43)

#2 Hash Join

on departments.location_id = locations.location_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.051ms | 48%

☰ Rows: 11 (Planned: 11)

\$ Cost: 0.13 (Total: 2.31)

#3 Seq Scan

on departments

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.004ms | 4%

☰ Rows: 11 (Planned: 11)

\$ Cost: 1.11 (Total: 1.11)

#4 Hash

Hash Node generates a hash table from the records in the input recordset. Hash is used by Hash Join.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.006ms | 6%

☰ Rows: 7 (Planned: 7)

#5 Seq Scan

on locations

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.035ms | 33%

☰ Rows: 7 (Planned: 7)

\$ Cost: 1.07 (Total: 1.07)

time	rows	estimation	cost
#1	Sort		
#2	↳ Nested Loop		
#3	↳ Hash Join		
#4	↳ Nested Loop		
#5	↳ Seq Scan		
#6	↳ Memoize		

#1 Sort

by jobs.job_title

Sort Node sorts a record set based on the specified sort key.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.544ms | 30%

Rows: 740 (Planned: 741)

Cost: 37.2 (Total: 107)

#2 Nested Loop

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.443ms | 24%

Rows: 740 (Planned: 741)

Cost: 24.5 (Total: 70)

#3 Hash Join

on employees.department_id = departments.department_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.258ms | 14%

Rows: 741 (Planned: 741)

Cost: 3.94 (Total: 45.5)

#4 Nested Loop

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.507ms | 28%

Rows: 741 (Planned: 741)

Cost: 22.9 (Total: 41.5)

#5 Seq Scan

on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.089ms | 5%

Rows: 741 (Planned: 741)

Cost: 18.4 (Total: 18.4)

#6 Memoize

Memoize Node is used to cache the results of the inner side of a nested loop. It avoids executing underlying nodes when the results for the current parameters are already in the cache.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0ms | 0%

Rows: ~741 (Planned: ~741)

Cost: 0.24 (Total: 0.24)

Loops: 741

time	rows	estimation	cost
#1	Sort		
#2	↳ Nested L...		
#3	↳ Hash J...		
#4	↳ Has...		
#5	↳ S...		
#6	↳ H...		

#1 Sort🕒 \$

by jobs.job_title

Sort Node sorts a record set based on the specified sort key.

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.524ms | 32%

☰ Rows: 740 (Planned: 741)

\$ Cost: 37.2 (Total: 87.8)

#2 Nested Loop🕒 \$

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.416ms | 26%

☰ Rows: 740 (Planned: 741)

\$ Cost: 24.5 (Total: 50.6)

#3 Hash Join🕒

on employees.department_id = departments.department_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.275ms | 17%

☰ Rows: 741 (Planned: 741)

\$ Cost: 3.93 (Total: 26.1)

#4 Hash Join🕒

on employees.job_id = jobs.job_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.298ms | 18%

☰ Rows: 741 (Planned: 741)

\$ Cost: 2.62 (Total: 22.2)

#5 Seq Scan\$

on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.102ms | 6%

☰ Rows: 741 (Planned: 741)

\$ Cost: 18.4 (Total: 18.4)

#6 Hash

Hash Node generates a hash table from the records in the input recordset. Hash is used by Hash Join.

General

IO & Buffers

Output

Workers

Misc

🕒 Timing: 0.01ms | 1%

☰ Rows: 19 (Planned: 19)

\$ Cost: 1.19 (Total: 1.19)

time	rows	estimation	cost
#1	Sort		
#2	↳ Nested L...		
#3	↳ Hash J...		
#4	↳ Has...		
#5	↳ S...		
#6	↳ H...		

#1 Sort⌚ \$

by jobs.job_title

Sort Node sorts a record set based on the specified sort key.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.491ms | 30%

≡ Rows: 740 (Planned: 741)

\$ Cost: 37.2 (Total: 87.8)

#2 Nested Loop⌚ \$

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.426ms | 26%

≡ Rows: 740 (Planned: 741)

\$ Cost: 24.5 (Total: 50.6)

#3 Hash Join⌚

on employees.department_id = departments.department_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.272ms | 17%

≡ Rows: 741 (Planned: 741)

\$ Cost: 3.93 (Total: 26.1)

#4 Hash Join⌚

on employees.job_id = jobs.job_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.316ms | 19%

≡ Rows: 741 (Planned: 741)

\$ Cost: 2.62 (Total: 22.2)

#5 Seq Scan\$

on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.105ms | 6%

≡ Rows: 741 (Planned: 741)

\$ Cost: 18.4 (Total: 18.4)

#6 Hash

Hash Node generates a hash table from the records in the input recordset. Hash is used by Hash Join.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.02ms | 1%

≡ Rows: 19 (Planned: 19)

\$ Cost: 1.19 (Total: 1.19)

time	rows	estimation	cost
#1	Sort		
#2	Nested Loop		
#3	Hash Join		
#4	Hash Join		
#5	Seq Scan		
#6	Hash		

#1 Sort

by jobs.job_title

Sort Node sorts a record set based on the specified sort key.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.558ms | 30%

Rows: 740 (Planned: 741)

Cost: 37.2 (Total: 86.7)

#2 Nested Loop

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.552ms | 29%

Rows: 740 (Planned: 741)

Cost: 24.4 (Total: 49.5)

#3 Hash Join

on employees.department_id = departments.department_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.243ms | 13%

Rows: 741 (Planned: 741)

Cost: 3.93 (Total: 25.1)

#4 Hash Join

on employees.job_id = jobs.job_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.422ms | 22%

Rows: 741 (Planned: 741)

Cost: 2.62 (Total: 21.2)

#5 Seq Scan

on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.106ms | 6%

Rows: 741 (Planned: 741)

Cost: 17.4 (Total: 17.4)

#6 Hash

Hash Node generates a hash table from the records in the input recordset. Hash is used by Hash Join.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.008ms | 0%

Rows: 19 (Planned: 19)

Cost: 1.19 (Total: 1.19)