

time	rows	estimation	cost
#1	HashAggr...		
#2	↳ Nested L...		
#3	↳ Seq Sc...		
#4	↳ Memo...		
#5	↳ Ind...		

#1 HashAggregate⌚\$🔊

by departments.department_name

HashAggregate Node groups records together based on a GROUP BY or aggregate function (like sum()). Hash Aggregate uses a hash to first organize the records by a key.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.284ms | 28%

☰ Rows: 11 (Planned: 200) | ⬆ over estimated by 18 ×

\$ Cost: 5.7 (Total: 45.4)

#2 Nested Loop⌚\$

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.454ms | 45%

☰ Rows: 741 (Planned: 741)

\$ Cost: 21.1 (Total: 39.7)

#3 Seq Scan⌚\$

on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.233ms | 23%

☰ Rows: 741 (Planned: 741)

\$ Cost: 18.4 (Total: 18.4)

#4 Memoize

Memoize Node is used to cache the results of the inner side of a nested loop. It avoids executing underlying nodes when the results for the current parameters are already in the cache.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0ms | 0%

☰ Rows: ~741 (Planned: ~741)

\$ Cost: 0.01 (Total: 0.24)

🔄 Loops: 741

#5 Index Scan

on departments

using departments_pkey

Index Scan Node finds relevant records based on an Index. Index Scans perform 2 read operations: one to read the index and another to read the actual value from the table.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.033ms | 3%

☰ Rows: ~11 (Planned: ~11)

\$ Cost: 0.23 (Total: 0.23)

🔄 Loops: 11