

time	rows	estimation	cost
#1	HashAggr...		
#2	└ Hash Join		
#3	└└ Seq Sc...		
#4	└└ Hash		
#5	└└└ Seq...		

#1 HashAggregate⌚ \$

by departments.department\_name

HashAggregate Node groups records together based on a GROUP BY or aggregate function (like sum()). Hash Aggregate uses a hash to first organize the records by a key.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.311ms | 33%

☰ Rows: 11 (Planned: 11)

\$ Cost: 3.81 (Total: 26.2)

#2 Hash Join⌚ \$

on employees.department\_id = departments.department\_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.452ms | 48%

☰ Rows: 741 (Planned: 741)

\$ Cost: 2.83 (Total: 22.4)

#3 Seq Scan⌚ \$

on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.169ms | 18%

☰ Rows: 741 (Planned: 741)

\$ Cost: 18.4 (Total: 18.4)

#4 Hash

Hash Node generates a hash table from the records in the input recordset. Hash is used by Hash Join.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.007ms | 1%

☰ Rows: 11 (Planned: 11)

#5 Seq Scan

on departments

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.009ms | 1%

☰ Rows: 11 (Planned: 11)

\$ Cost: 1.11 (Total: 1.11)