

time	rows	estimation	cost
#1	HashAggr...		
#2	└ Hash Join		
#3	└└ Seq Sc...		
#4	└└ Hash		
#5	└└└ Seq...		

#1 HashAggregate⌚

by locations.city

HashAggregate Node groups records together based on a GROUP BY or aggregate function (like sum()). Hash Aggregate uses a hash to first organize the records by a key.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.011ms | 16%

≡ Rows: 7 (Planned: 11)

\$ Cost: 0.16 (Total: 14.8)

#2 Hash Join⌚

on locations.location_id = departments.location_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.015ms | 22%

≡ Rows: 11 (Planned: 11)

\$ Cost: 1.15 (Total: 14.7)

#3 Seq Scan⌚💰🗨

on locations

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.007ms | 10%

≡ Rows: 7 (Planned: 240) | ⬆ over estimated by 34 ×

\$ Cost: 12.4 (Total: 12.4)

#4 Hash⌚

Hash Node generates a hash table from the records in the input recordset. Hash is used by Hash Join.

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.009ms | 13%

≡ Rows: 11 (Planned: 11)

#5 Seq Scan⌚

on departments

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General

IO & Buffers

Output

Workers

Misc

⌚ Timing: 0.026ms | 38%

≡ Rows: 11 (Planned: 11)

\$ Cost: 1.11 (Total: 1.11)