Plan Raw Query Stats 1.5.0

Execution time: N/A | Planning time: 0.589ms | Triggers: N/A | Settings

time	rows	estimation	c
#1 HashAggr...			
#2 └ Hash Join			
#3 ├ Seq Sc...			
#4 └ Hash			
#5 └ Seq...			

#1 HashAggregate
by `locations.city`

HashAggregate Node groups records together based on a GROUP BY or aggregate function (like sum()). Hash Aggregate uses a hash to first organize the records by a key.

General | IO & Buffers | Output | Workers | Misc

Timing: 0.01ms | 9%
Rows: 7 (Planned: 7)
Cost: 0.12 (Total: 2.43)

#2 Hash Join
on `departments.location_id = locations.location_id`

Hash Join Node joins two record sets by hashing one of them (using a **Hash Scan**).

General | IO & Buffers | Output | Workers | Misc

Timing: 0.051ms | 48%
Rows: 11 (Planned: 11)
Cost: 0.13 (Total: 2.31)

#3 Seq Scan
on `departments`

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General | IO & Buffers | Output | Workers | Misc

Timing: 0.004ms | 4%
Rows: 11 (Planned: 11)
Cost: 1.11 (Total: 1.11)

#4 Hash

Hash Node generates a hash table from the records in the input recordset. Hash is used by **Hash Join**.

General | IO & Buffers | Output | Workers | Misc

Timing: 0.006ms | 6%
Rows: 7 (Planned: 7)

#5 Seq Scan
on `locations`

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

General | IO & Buffers | Output | Workers | Misc

Timing: 0.035ms | 33%
Rows: 7 (Planned: 7)
Cost: 1.07 (Total: 1.07)