

time	rows	estimation	cost
#1	Sort		
#2	↳ Nested L...		
#3	↳ Hash J...		
#4	↳ Nes...		
#5	↳ S...		
#6	↳ ...		

#1 Sort🕒💰
by jobs.job_title

Sort Node sorts a record set based on the specified sort key.

GeneralIO & BuffersOutputWorkersMisc

🕒 Timing: 0.544ms | 30%
☰ Rows: 740 (Planned: 741)
💰 Cost: 37.2 (Total: 107)

#2 Nested Loop🕒💰

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

GeneralIO & BuffersOutputWorkersMisc

🕒 Timing: 0.443ms | 24%
☰ Rows: 740 (Planned: 741)
💰 Cost: 24.5 (Total: 70)

#3 Hash Join🕒

on employees.department_id = departments.department_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

GeneralIO & BuffersOutputWorkersMisc

🕒 Timing: 0.258ms | 14%
☰ Rows: 741 (Planned: 741)
💰 Cost: 3.94 (Total: 45.5)

#4 Nested Loop🕒💰

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

GeneralIO & BuffersOutputWorkersMisc

🕒 Timing: 0.507ms | 28%
☰ Rows: 741 (Planned: 741)
💰 Cost: 22.9 (Total: 41.5)

#5 Seq Scan💰
on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

GeneralIO & BuffersOutputWorkersMisc

🕒 Timing: 0.089ms | 5%
☰ Rows: 741 (Planned: 741)
💰 Cost: 18.4 (Total: 18.4)

#6 Memoize

Memoize Node is used to cache the results of the inner side of a nested loop. It avoids executing underlying nodes when the results for the current parameters are already in the cache.

GeneralIO & BuffersOutputWorkersMisc

🕒 Timing: 0ms | 0%
☰ Rows: ~741 (Planned: ~741)
💰 Cost: 0.24 (Total: 0.24)
🔄 Loops: 741