

time	rows	estimation	cost
#1	Sort		
#2	Nested Loop		
#3	Hash Join		
#4	Hash Join		
#5	Seq Scan		
#6	Hash		

#1 Sort

by jobs.job_title

Sort Node sorts a record set based on the specified sort key.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.558ms | 30%

Rows: 740 (Planned: 741)

Cost: 37.2 (Total: 86.7)

#2 Nested Loop

Nested Loop Node merges two record sets by looping through every record in the first set and trying to find a match in the second set. All matching records are returned.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.552ms | 29%

Rows: 740 (Planned: 741)

Cost: 24.4 (Total: 49.5)

#3 Hash Join

on employees.department_id = departments.department_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.243ms | 13%

Rows: 741 (Planned: 741)

Cost: 3.93 (Total: 25.1)

#4 Hash Join

on employees.job_id = jobs.job_id

Hash Join Node joins two record sets by hashing one of them (using a Hash Scan).

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.422ms | 22%

Rows: 741 (Planned: 741)

Cost: 2.62 (Total: 21.2)

#5 Seq Scan

on employees

Seq Scan Node finds relevant records by sequentially scanning the input record set. When reading from a table, Seq Scans (unlike Index Scans) perform a single read operation (only the table is read).

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.106ms | 6%

Rows: 741 (Planned: 741)

Cost: 17.4 (Total: 17.4)

#6 Hash

Hash Node generates a hash table from the records in the input recordset. Hash is used by Hash Join.

GeneralIO & BuffersOutputWorkersMisc

Timing: 0.008ms | 0%

Rows: 19 (Planned: 19)

Cost: 1.19 (Total: 1.19)